

Day 09 - Logistic Regression

Oct. 6, 2020



Administrative

- **Homework 3** will be assigned Friday 10/9 and due Friday 10/23
- **Midterm** will be given Thursday 10/29 in class

From Pre-Class Assignment

Useful Stuff

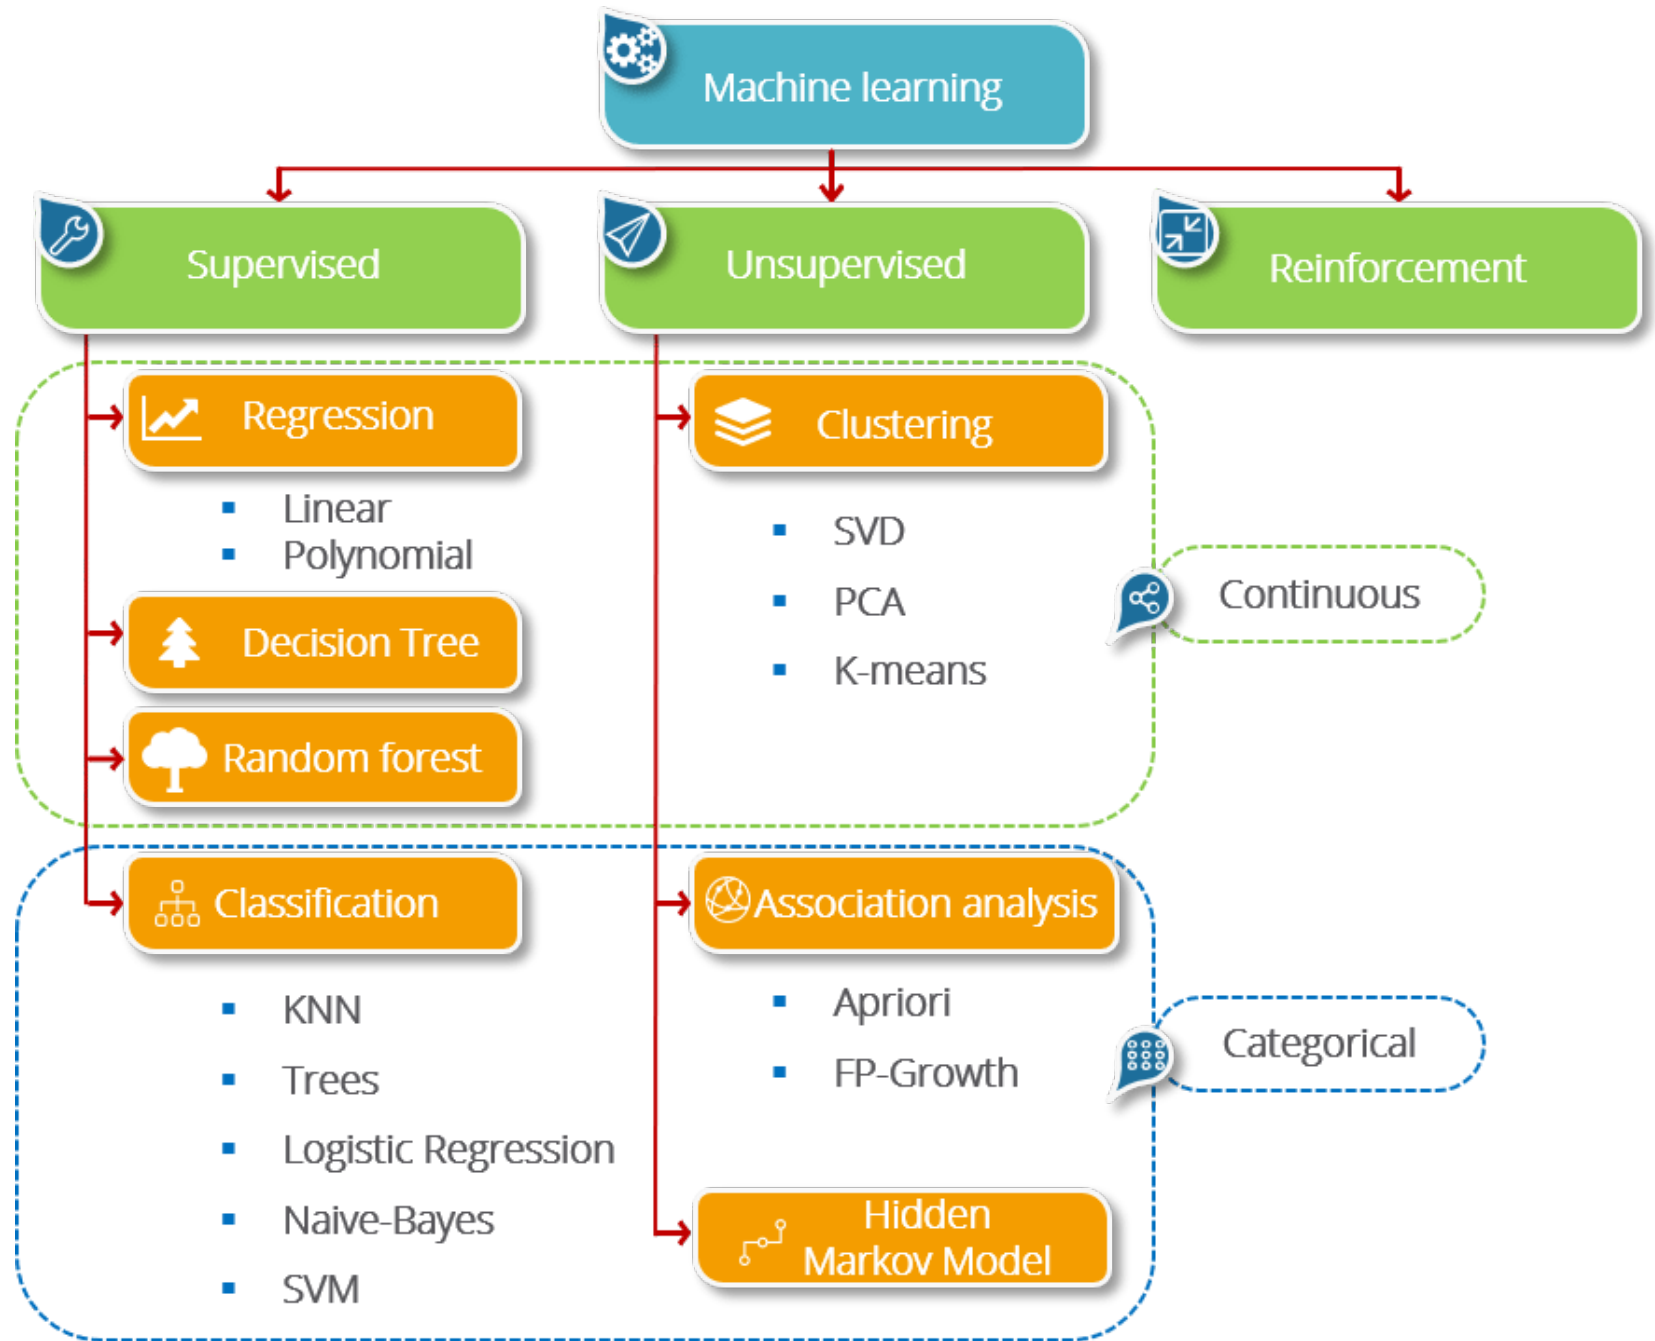
- Videos from Google were helpful to understand the scope of Machine Learning
- I have a better understanding of train/test split

Challenging bits

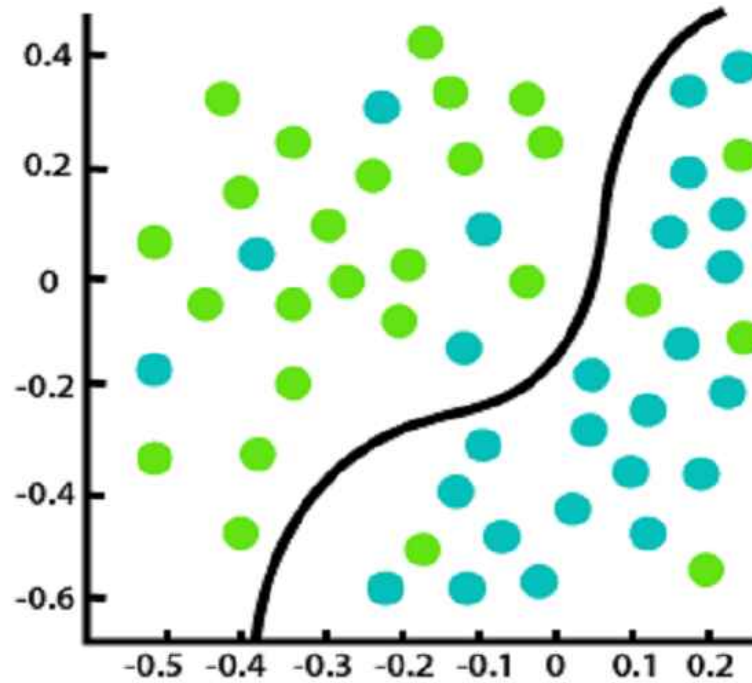
- I am still a little confused about why we split the data
- I am not sure what `make_classification` is doing
- What are redundant and informative features? How do we see them in the plots?

We will be doing classification tasks for a few weeks, so we will get lots of practice

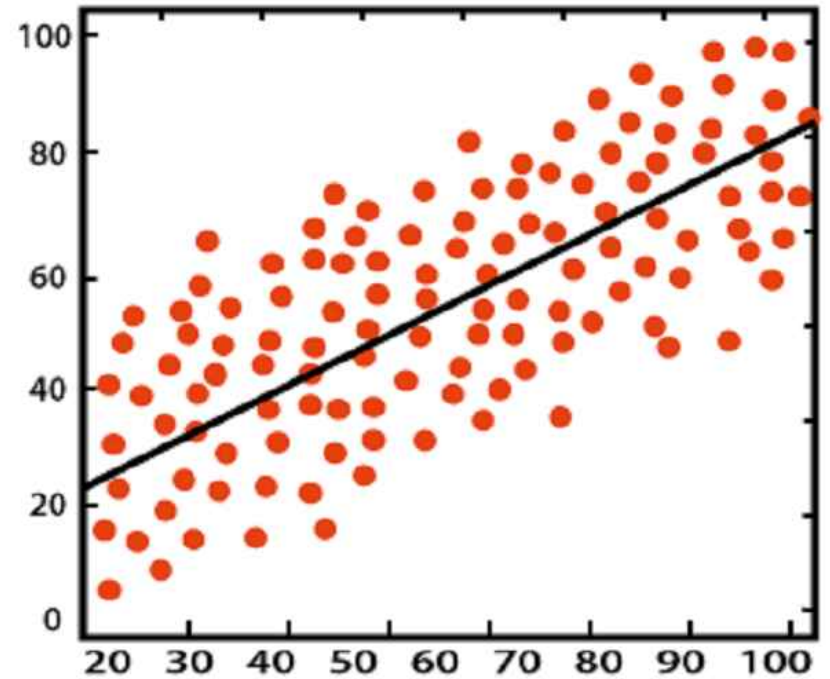
Machine Learning



Classification



Classification



Regression

Classification Algorithms

- **Logistic Regression:** The most traditional technique; was developed and used prior to ML; fits data to a "sigmoidal" (s-shaped) curve; fit coefficients are interpretable
- **K Nearest Neighbors (KNN):** A more intuitive method; nearby points are part of the same class; fits can have complex shapes
- **Support Vector Machines (SVM):** Developed for linear separation (i.e., find the optimal "line" to separate classes; can be extended to curved lines through different "kernels")
- **Decision Trees:** Uses binary (yes/no) questions about the features to fit classes; can be used with numerical and categorical input
- **Random Forest:** A collection of randomized decision trees; less prone to overfitting than decision trees; can rank importance of features for prediction
- **Gradient Boosted Trees:** An even more robust tree-based algorithm

We will learn Logistic Regression, KNN, and SVM, but `sklearn` provides access to the other three methods as well.

Generate some data

`make_classification` lets us make fake data and control the kind of data we get.

- `n_features` - the total number of features that can be used in the model
- `n_informative` - the total number of features that provide unique information for classes
 - say 2, so x_0 and x_1
- `n_redundant` - the total number of features that are built from informative features (i.e., have redundant information)
 - say 1, so $x_2 = c_0 x_0 + c_1 x_1$
- `n_class` - the number of class labels (default 2: 0/1)
- `n_clusters_per_class` - the number of clusters per class

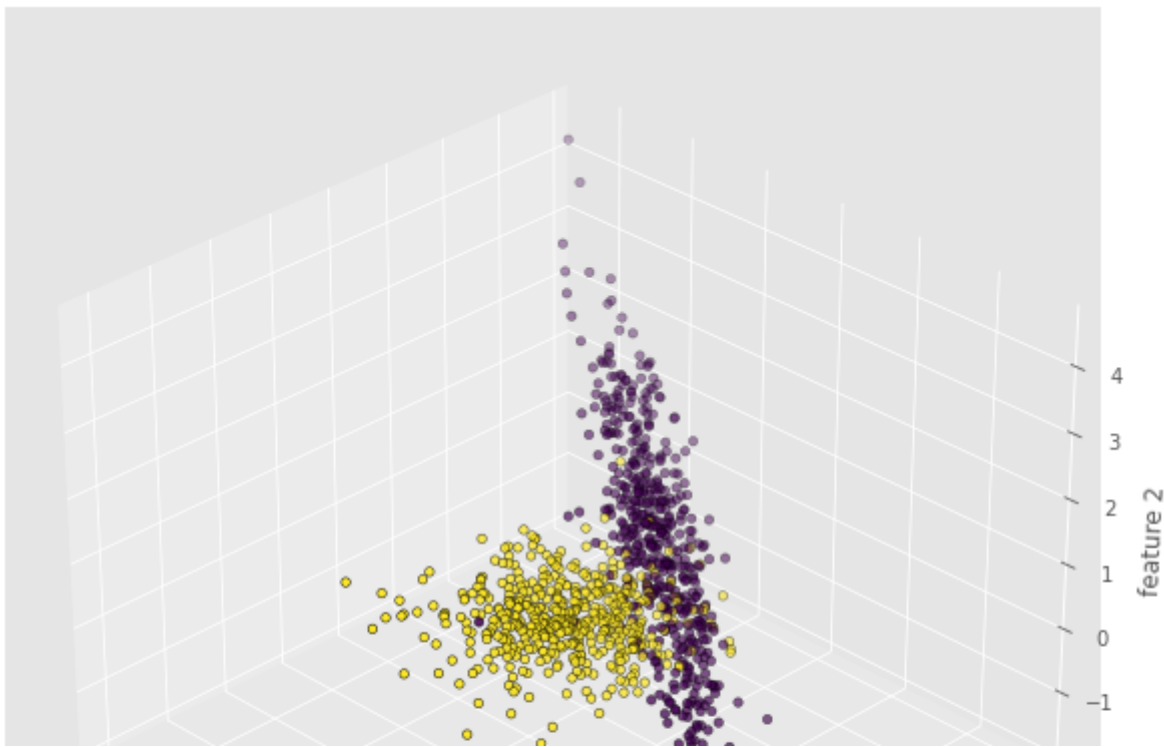
[illegible]

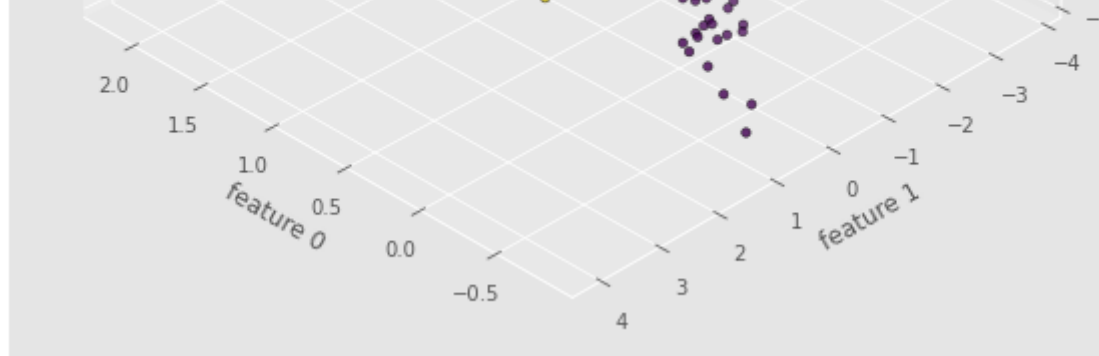

```
In [64]: ## Let's look at these 3D data
from mpl_toolkits.mplot3d import Axes3D
fig = plt.figure(figsize=(8,8))
ax = Axes3D(fig, rect=[0, 0, .95, 1], elev=30, azimuth=135)

xs = features[:, 0]
ys = features[:, 1]
zs = features[:, 2]

ax.scatter3D(xs, ys, zs, c=class_labels, ec='k')
ax.set_xlabel('feature 0')
ax.set_ylabel('feature 1')
ax.set_zlabel('feature 2')
```

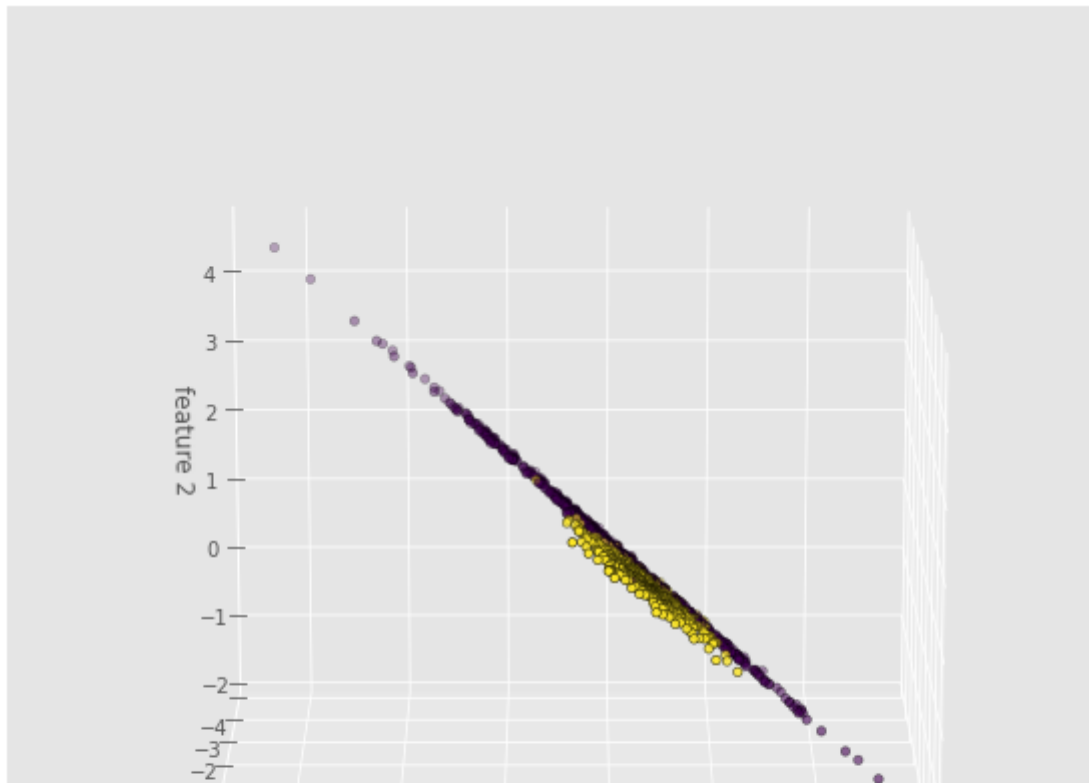
```
Out[64]: Text(0.5, 0, 'feature 2')
```

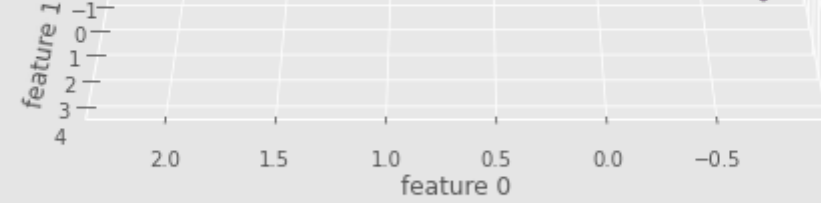




```
In [65]: ## From a different angle, we see the 2D nature of the data  
fig = plt.figure(figsize=(8,8))  
ax = Axes3D(fig, rect=[0, 0, .95, 1], elev=15, azimuth=90)  
  
xs = features[:, 0]  
ys = features[:, 1]  
zs = features[:, 2]  
  
ax.scatter3D(xs, ys, zs, c=class_labels, ec = 'k')  
ax.set_xlabel('feature 0')  
ax.set_ylabel('feature 1')  
ax.set_zlabel('feature 2')
```

```
Out[65]: Text(0.5, 0, 'feature 2')
```





Feature Subspaces

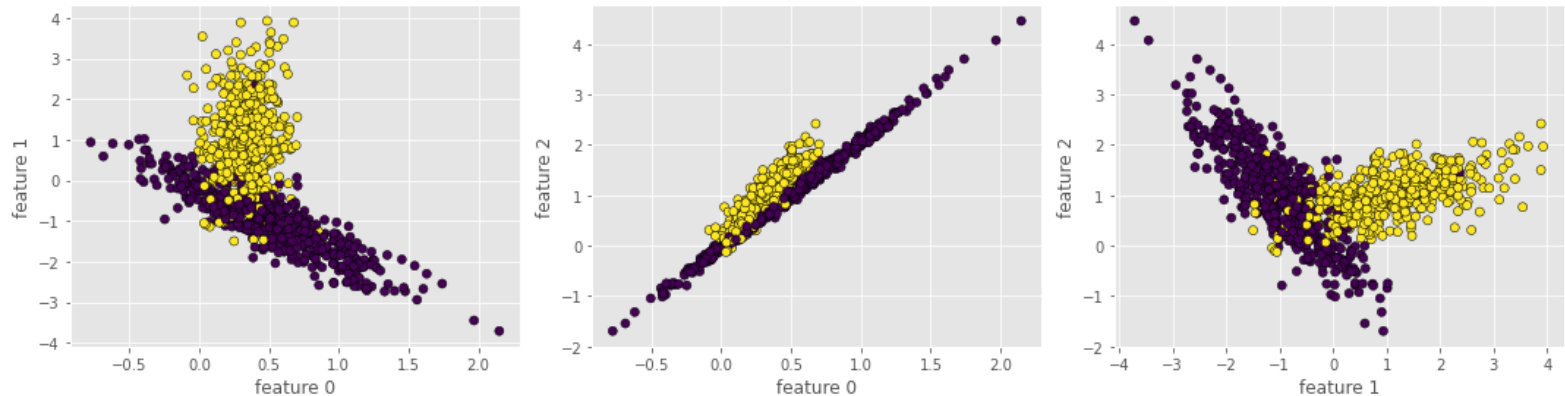
For higher dimensions, we have take 2D slices of the data (called "projections" or "subspaces")

```
In [66]: f, axs = plt.subplots(1,3,figsize=(15,4))
plt.subplot(131)
plt.scatter(features[:, 0], features[:, 1], marker = 'o', c = class_labels, ec =
'k')
plt.xlabel('feature 0')
plt.ylabel('feature 1')

plt.subplot(132)
plt.scatter(features[:, 0], features[:, 2], marker = 'o', c = class_labels, ec =
'k')
plt.xlabel('feature 0')
plt.ylabel('feature 2')

plt.subplot(133)
plt.scatter(features[:, 1], features[:, 2], marker = 'o', c = class_labels, ec =
'k')
plt.xlabel('feature 1')
plt.ylabel('feature 2')

plt.tight_layout()
```



What about Logistic Regression?

Logistic Regression attempts to fit a sigmoid (S-shaped) function to your data. This shape assumes that the probability of finding class 0 versus class 1 increases as the feature changes value.

```

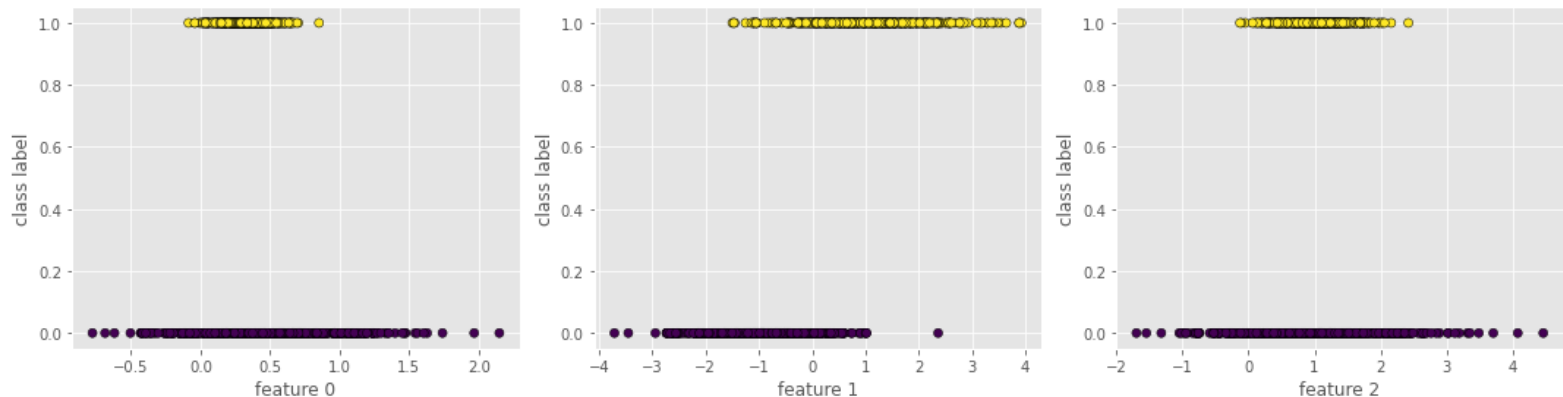
In [70]: f, axs = plt.subplots(1,3,figsize=(15,4))
plt.subplot(131)
plt.scatter(features[:,0], class_labels, c=class_labels, ec='k')
plt.xlabel('feature 0')
plt.ylabel('class label')

plt.subplot(132)
plt.scatter(features[:,1], class_labels, c=class_labels, ec='k')
plt.xlabel('feature 1')
plt.ylabel('class label')

plt.subplot(133)
plt.scatter(features[:,2], class_labels, c=class_labels, ec='k')
plt.xlabel('feature 2')
plt.ylabel('class label')

plt.tight_layout()

```



Questions, Comments, Concerns?